

Fast XML Parsing with Haskell

Neil Mitchell

<http://ndmitchell.com>

+ Christopher Done




CONTENT DISCLAIMER

Optimisation is the art of making something faster

- Desire: It must go too slow
- Benchmark: You must know how fast it goes
- Profile: You must know what to change

System Optimisation

- Optimisation folk lore
 - 90% of the time is spent running 100 lines
 - Optimise those 100 lines and profit



Warning: After a few rounds of optimisation, your profile may be mostly flat

Haskell inner loops

C

Security!!!!
Painful allocation
Marshalling
No abstractions
Single lump
Less familiar
Verbose
Undefined behaviour
Portability
Segfaults

Haskell

Security!
Implicit allocation
INLINE and -O2
Many abstractions

The Problem

- Parse XML to a DOM tree and query it for tags/attributes

```
<conference title="Haskell eXchange" year=2017>  
  <talk author="Julie Moronuki">  
    A Monoid for All Seasons  
  </talk>  
  <talk author="Neil Mitchell">  
    Fast XML parsing with Haskell  
    <active/> <!-- remove this in 30 mins -->  
  </talk>  
</conference>
```

Existing Solutions

- xml – 100x-300x slower
- expat – 40x-100x slower
- xml-conduit – much slower
- tagsoup – SAX based
- XMLParser
- xmlhtml
- xml-pipe
- PugiXML: C++ library, fastest by a lot
 - Haskell binding segfaults ☹️

PugiXML Tricks

- Extremely fast – faster than all others
 - 9x faster than libxml
 - 27x faster than msxml
 - Closest are asxml (x86 only), rapidxml
 - “Parsing XML at the Speed of Light”
- Ignore the DOCTYPE stuff (no one cares)
- Does not validate
- In-place parsing

Our Tricks


- Ignore the DOCTYPE stuff (no one cares)
- Does not validate
- In-place parsing (even more so)
- Don't expand entities e.g. &
 - All returned strings are offsets into the source
 - In body text, only care about >, so memchr
- Hexml: Haskell friendly C library + wrapper
- Xeno: Pure Haskell alternative

Approach 1: C inner loops

Hexml

<https://hackage.haskell.org/package/hexml>

Hexml Memory



Hexml Interface (types)

```
typedef struct
{
    int32_t start;
    int32_t length;
} str;
```

```
typedef struct
{
    str name; // tag name, e.g. <[foo]>
    str inner; // inner text, <foo>[bar]</foo>
    str outer; // outer text, [<foo>bar</foo>]
} node;
```

Hexml Interface (functions)

```
document* document_parse(const char* s, int slen);
```

```
char* document_error(const document* d);
```

```
void document_free(document* d);
```

```
node* document_node(const document* d);
```

```
attr* node_attributes(const document* d, const node* n, int* res);
```

```
attr* node_attribute(const document* d, const node* n, const char* s,  
                    int slen);
```

How did I get to that?

- I've written FFI bindings before, so know what is hard/slow, and avoided it!
 - Simple memory management (only document)
 - Functions are relatively big – where possible known structs are used
 - Use ByteString because it is FFI friendly (C ptr)
- Intuition and experience matters...
 - (My excuse for not using a simple example)

Wrapping Haskell (types)

```
data Str = Str {  
    strStart :: Int32,  
    strLength :: Int32  
}
```

```
instance Storable Str where
```

```
    sizeof _ = 8
```

```
    alignment _ = alignment (0 :: Int64)
```

```
    peek p = Str <$> peekByteOff p 0 <*> peekByteOff p 4
```

```
    poke p (Str a b) = pokeByteOff p 0 a >> pokeByteOff p 4 b
```

```
typedef struct  
{  
    int32_t start;  
    int32_t length;  
} str;
```

Wrapping Haskell (functions)

```
document* document_parse(const char* s, int slen);  
void document_free(document* d);  
node* document_node(const document* d);
```

data CDocument

data CNode

foreign import ccall document_parse
:: CString -> CInt -> IO (Ptr CDocument)

foreign import ccall "&document_free" document_free
:: FunPtr (Ptr CDocument -> IO ())

foreign import ccall unsafe document_node
:: Ptr CDocument -> IO (Ptr CNode)

Wrapping Haskell (memory)

- Document is not on the Haskell API (pretend it's a node)
- A node must know about the text of it, the document it is in, and the node itself

```
data Node = Node  
  BS.ByteString  
  (ForeignPtr CDocument)  
  (Ptr CNode)
```


Creating Node

```
parse :: BS.ByteString -> BS.ByteString Node
parse src = unsafePerformIO $
  BS.unsafeUseAsCStringLen src $ \(str, len) -> do
    doc <- document_parse str (fromIntegral len)
    doc <- newForeignPtr document_free doc
    node <- document_node doc
    return $ Node src doc node
```

Using Node

```
attr* node_attributes(const document* d, const node* n, int* res);  
node_attributes :: Ptr CDocument -> Ptr CNode -> Ptr CInt -> IO (Ptr CAttr)
```

```
attributes :: Node -> [Attribute]
```

```
attributes (Node src doc n) = unsafePerformIO $
```

```
  withForeignPtr doc $ \d ->
```

```
    alloca $ \count -> do
```

```
      res <- node_attributes d n count
```

```
      count <- fromIntegral <$> peek count
```

```
      return [attrPeek src doc $ plusPtr res $ i*szAttr  
              | i <- [0..count-1]]
```

The big picture

- Define some simple functions types in C
 - Wrap them to Haskell almost mechanically
- Define some types in C
 - Wrap them to Haskell in a context specific way
- Wrap the functions into usable Haskell
 - Requires smarts to get them looking right
 - Requires insane attention to detail to not segfault
- Note we *haven't* shown the C code!

Continuing onwards

- Testing can and should be in Haskell
 - Explicit test cases based on errors
 - Property based testing
 - Wrote a renderer, checked for idempotence
- Debugging C by printf is super painful
 - I used Visual Studio for interactive debugging
 - Used American Fuzzy Lop for fuzzing

Results

- Fast! ~2x faster than PugiXML
- Simple! Nice clean interface
- Abstractable! hexml-lens has lenses on top
- But ran into...
 - Undefined behaviour in C
 - Buffer overruns in C
 - Incorrect memory usage in Haskell
- All removed with blood, sweat and tears

Approach 2: Haskell inner loops

Xeno

<https://hackage.haskell.org/package/xeno>

Christopher Done, now Marco Zocca

Approach

- Hexml: Think hard and be perfect
- Xeno: Follow this methodology
 - Watch memory allocations like a hawk
 - Start simple, benchmark
 - Add features, rebenchmark
 - Build from composable pieces

Simplest possible

`parseTags :: ByteString -> Int -> ()` -- walk a document

`parseTags str |`

`| Just i <- findNext '<' str |`

`, Just i <- findNext '>' str (i+1)`

`= parseTags str (i+1)`

`| otherwise = ()`

`findNext :: Char -> ByteString -> Int -> Maybe Int`

`{-# INLINE findNext #-}`

`findNext c str offset = (+ offset) <$>`

`BS.elemIndex c (BS.drop offset str)`

Timing

File	hexml	xeno
4KB	6.395 μ s	2.630 μ s
42KB	37.55 μ s	7.814 μ s

- Basically measuring C memchr function
 - Plus bounds checking!
- Shows Haskell is not adding huge overhead

<https://hackage.haskell.org/package/criterion>

Memory

Case	Bytes	GCs	Check
4kb parse	1,168	0	OK
42kb parse	1,560	0	OK
52kb parse	1,168	0	OK
182kb parse	1,168	0	OK

- Memory usage is linear – not per $\langle \rangle$ pair
- Don't we allocate a Just per $\langle \rangle$?

<https://hackage.haskell.org/package/weigh>

Watching the Just

```
parseTags str i  
  | Just i <- findNext '<' str i
```

```
{-# INLINE findNext #-}  
findNext c str offset = (+ offset) <$>  
  BS.elemIndex c (BS.drop offset str)
```

```
{-# INLINE elemIndex #-}  
BS.elemIndex str x =  
  let q = memchr str x  
  in if q == nullPtr then Nothing else Just $ str - q
```

Is 'Just' expensive?

- A single Just requires:
 - Heap check (comparison, one per function)
 - Alloc (addition)
 - Construction (memory writes)
 - Examination (memory reads, jump)
 - GC (expensive, one every so often)
- Not “expensive”, just not free

Incrementally add bits

- Parse comments, tags, attributes
- Return results
- At each step:
 - Benchmark (will slow down a bit)
 - Memory (should remain zero)
- Tricks
 - INLINE, -O2, alternative functions

Making it useful

parseTags

```
:: (s -> ByteString -> s)
-> ByteString -> Int -> s
-> Either XenoException s
```

parseTags fTag str l s

```
| Just i <- findNext '<' str l = case findNext '>' str (i+1) of
  Nothing -> Left $ XenoParseError "mismatched <"
  Just j -> parseTags fTag str (i+1) $ fTag s $ BS.substr (i+1) j
| otherwise = Right s
```

*Xeno specialises to a Monad and uses impure exceptions.
Does that make it go faster or slower?*

SAX Parser

fold

```
:: (s -> ByteString -> s) -- ^ Open tag.  
-> (s -> ByteString -> ByteString -> s) -- ^ Attribute.  
-> (s -> ByteString -> s) -- ^ End of open tag.  
-> (s -> ByteString -> s) -- ^ Text.  
-> (s -> ByteString -> s) -- ^ Close tag.  
-> s  
-> ByteString  
-> Either XenoException s
```

DOM Parser

- Can be built on top of the SAX parser
 - Beautiful abstraction in action
- Harder problem
 - Can't aim for zero allocations
 - Need a smart compact data structure
 - Need ST, STURef, vector

Xeno vs Hexml

File	hexml-dom	xeno-sax	xeno-dom
4KB	6.123 μ s	5.038 μ s	10.35 μ s
31KB	9.417 μ s	2.875 μ s	5.714 μ s
211KB	256.3 μ s	240.4 μ s	514.2 μ s

Haskell inner loops

C

Security!!!!

Painful allocation

Marshalling

No abstractions

Single lump

Less familiar

Verbose

Undefined behaviour

Portability

Segfaults

Haskell

Security!

Implicit allocation

Many abstractions

INLINE and -O2

GHC version tuning

Slower

Ongoing compromise

Conclusion

- C is up front design, Haskell is feedback
- Haskell can use better abstraction and security
- C is a lot harder than I remember
- Haskell FFI is exceptionally good

I *personally* prefer C inner loops to Haskell

DOM Storage

- Now onto a smart representation/algo
 - Haskell and C share the same ideas
 - C inner loops requires DOM storage also in C
- Needs to be compact
 - Store attributes and nodes in single alloc
- Easier to describe in C?

DOM Attributes

typedef struct

{

int size; // number used

int used; // number available, doubles

attr* attrs; // dynamically allocated buffer

attr* alloc; // what to call free on

} attr_buffer;

Buffer that doubles on reallocation

Plus fast path for special allocation

DOM Document

typedef struct

```
{  
    const char* body; // pointer to initial argument  
                        // not owned by us  
  
    char* error_message;  
    node_buffer nodes;  
    attr_buffer attrs;  
} document;
```

Nothing interesting

DOM Creation

```
typedef struct  
{  
    document document;  
    attr attrs[1000];  
    node nodes[500];  
} buffer;
```

*Alloc a buffer, point document.nodes at buffer.nodes
If resizing, just ignore the memory
1 allocation for 3 buffers*

DOM Nodes

```
typedef struct
```

```
{
```

```
    int size;
```

```
    int used_front; // front entries, stored for good
```

```
    int used_back; // back entries, stack based, copied into front
```

```
    node* nodes; // dynamically allocated buffer
```

```
    node* alloc; // what to call free on
```

```
} node_buffer;
```

Want all DOM children to be adjacent (compact)

What about nested children?

Copy to the end of the buffer, then commit

Resizing needs to copy too



C is hard: [1/7]

```
static inline bool is(char c, char tag)
{
    return table[(unsigned char) c] & tag;
}
```

PORTABILITY

OUT OF BOUNDS READ



C is hard: [2/7]

```
if (get peek(d) != '=')  
{  
    set_error(d, "Expected = in attribute, but  
missing");  
    return start_length(0, 0);  
}  
skip(d, 1);
```

INCORRECT RESULT



C is hard: [3/7]

```
attributeBy (Node src doc n) str =  
  unsafePerformIO $ withForeignPtr doc $ \d ->  
    BS.unsafeUseAsCStringLen str $ \(bs, len) -> do  
      r <- node_attributeBy d n bs $ fromIntegral len  
      touchForeignPtr $ fst3 $ BS.toForeignPtr src  
      return $ if r == nullPtr then Nothing  
               else Just $ attrPeek src doc r
```

USE AFTER FREE



C is hard: [4/7]

```
let src0 = src <> BS.singleton '\0'  
...  
return $ Node src0 doc node
```

USE AFTER FREE



C is hard: [5/7]

~~`d->nodes.nodes[0].nodes = parse_content(d);`~~

`str content = parse_content(d);`
`d->nodes.nodes[0].nodes = content;`

UNDEFINED
BEHAVIOUR

UNPORTABLE

USE AFTER FREE



C is hard: [6/7]

```
if (peek_at(d, -3) == '-' &&  
    peek_at(d, -2) == '-')
```

INCORRECT RESULT



C is hard: [7/7]

```
while (1 d->error_message == NULL)
```

```
if (d->error_message != NULL) return;
```

```
c = get(d);
```

OUT OF BOUNDS READ