



Pyrefly

A Python typechecker and language server

Neil Mitchell, Meta



Python has types!

```
def test(x: int) -> str:  
    return f"test{x}"
```

- Since version 3.5 (2015) in PEP 484
- Not checked by default - the interpreter skips* them
 - Runtime checking - typeguard, pydantic, enforce - runtime cost
 - Static checking - mypy, pyright, pyre, pytype, pyrefly, ty, zuban
- More complex than you might think!
 - Generics, literals, higher-order, dataclass transforms, narrowing
 - A standards document, conformance tests - constantly evolving



What is Pyrefly?

- A static checker and language server (IDE/LSP provider)
- An open-source standards-compliant Python type checker
- Fast and parallel (written in Rust)
- From the team behind Pyre (no code in common)
- Lots of type inference

pyrefly.org



Sandbox (pyrefly.org/sandbox)

```
1  from typing import *
2
3  def test(x: int) -> str:
4      return f"{x}"
5
6  y: list[str] = []
7  y.append(test(42))
8  test(y[0]).
```

ERROR 8:6-10: Argument `str` is not assignable to parameter
`x` with type `int` in function `test` [bad-argument-type]

capitalise

BoundMethod[str, Overload[(self: Lit...]

casefold

center

count

encode



History of Pyrefly

- Meta develops Instagram which is a massive codebase of Python
 - Over 20 million lines of Python
 - 3B monthly active users
 - Over 3.3K daily Python developers across Meta
- In 2017 we started work on Pyre
 - MyPy was considered, but found to be too slow
 - Descendent of Hack (PHP) and Flow (Javascript)
 - Written in OCaml
 - Essential to our development flow



The problems with Pyre

OCaml wasn't a great choice at Meta

- Didn't work on Windows
- Parallelism was hard (multiprocess)
- Barrier to open source contributors
- Some problems were solved later, outside Meta




But Pyre had issues too

- Started as abstract interpretation (fixed points, desugaring)
- Started as a command line, hard to pivot to IDE
- Started closed source, never focused on open source users



The birth of Pyrefly

- August 2024 two of us started prototyping MiniPyre
 - 7 prototypes written, constraints, subset based, abstract interpretation...
 - Using Rust (cross platform and fast)
 - Hard bits first: generics, recursion, overloads, `import *`
- October 2024 it was working well, so we started ~~Pyre2~~ Pyrefly
 - Implement features, following the typing spec
 - Implement LSP
- May 2025 we released an alpha at PyConUS
- October 2025 we are still improving, but give it a go (pyrefly.org)





The Python type system



Why types?

- Python developers are buying in more
 - “2025 is the year of type-checking for Python. I’m so excited.”
- Faster inner loop - run the code less
- Spot typos
- Make corner cases safer
- Understand the code better, documentation, goto-def
- LLM grader
- Write code faster - auto-completion
- Machine checked documentation
- Refactor with peace of mind



The aim

- Give a way to annotate existing code and detect bugs
- The trend has been increasingly complex type system features to model the code people actually write
- Typing conformance test (Pyrefly gets 67.3%, similar to mypy)
- Detailed specification (<https://typing.python.org/>)
 - But still lots of choices to make as to precise semantics
 - Inference left unspecified



Basic types

- A few places you can write types
- The basic primitives, plus class types

```
class MyType:  
    field: bool  
    pass  
  
def test(x: int) -> str:  
    value: MyType = ...  
    return f"test{x}"
```



Literals

- For int, str, bool, bytes, you can write literals that restrict the type

```
def open_file(mode: Literal['r', 'w']): ...  
  
def calc() -> Literal[42]:  
    return 42  
  
x: Literal['test'] = 'te' + 'st'
```



Aliases and forward refs

- You can define type aliases as values
- Types in strings are OK, to deal with execution order
- Is `z = "hello"` a type?

```
def f(v: "X") -> "Y":  
    return v + 1  
  
X = int  
  
print(X("42"))  
  
type Y = "int"
```



Union and gradual types

Union[A, B] == A | B

Any ~= int | bool | MyType | ...

```
def test(x: Any | bool):
```

```
    ...
```

- Any represents an unknown static type
- Every type is assignable to Any, and Any is assignable to every type.
- A function parameter without an annotation is assumed to be Any.
- Also Never/NoReturn for the empty union.

Is an unannotated variable implicitly annotated with Any? Systems vary.



Generics

Two forms, using TypeVar and generic syntax

TypeVar can specify variance,
generic syntax infers it

Can specify constraints on both,
e.g. must be iterable.

```
X = TypeVar("X")
def box1(x: X) -> list[X]:
    return [x]

def box2[Y](y: Y) -> list[Y]:
    return [y]
```



Protocols

Structural subtyping

Standard types for iteration,
collections, mutable collections,
context managers etc

```
class SupportsClose(Protocol):  
    def close(self) -> None: ...
```



Overloads

A set of “fake signatures”, where only one matches.

```
@overload
def not(x: Literal[True]) -> Literal[False]: ...
@overload
def not(x: Literal[False]) -> Literal[True]: ...
def not(x: Any) -> bool:
    return False if x else True
```



Callable

Specify functions as types. Versatile with generics and concat.
But awkward, since doesn't let you specify argument names.

```
type Simple = Callable[[int, str], bool]
def f[**P, T](
    call: Callable[Concatenate[int, P], T],
    *args: P.args, **kwargs: P.kwargs) -> T:
    return call(42, *args, **kwargs)
```



Narrowing

Flow control refines the types in branches

```
def f(x: str | None):  
    if x:  
        return x.capitalize()  
    else:  
        return "none"
```



Data class transforms

PEP-681, Dec 2021 - SQL/ORM style solution





Typeshed library

- Attempts to give types to everything in the standard library.
- A little awkward - the standard library was written without types.
- Often there is what the library does precisely, and what you would have made it do if you knew about types.
 - E.g. protocols are close, rather than precise.



The Pyrefly design



Design goals

- Efficient in time and memory. Incremental. Must cope with Instagram.
- Flexible. Command line. IDE/LSP. Buck. MCP. Dune?
- Hackable. Solid principles. Simple code.
- Good user experience.
- Deal with circular import graphs.

Conclusion*: Build system, operate at the *file* level, evict old data.

Implication: Each file must be super quick (a few ms)

* Caveat: I describe everything in my life as a build system



Architecture



Build
System



1. Code

- Read the code off disk (or if LSP, from LSP messages).



2. AST

- Parse the code into an AST.
- Uses the Ruff parser from Astral (which is amazing)
 - Always succeeds, error correcting parser
 - Which is like the ultimate fuzzer...



3. Exports

- What symbols does each module export
- Modules always export all their imports, apart from builtins
- Not trivial because of import * - can require a fixed point
- Required to know which module provides a given value

```
from module1 import *
from module2 import *
from module3 import y

x = z
```



4. Bindings (statements)

- How do statements relate to each other, where to variables flow
- Key -> Value mapping, where Value's contain other Key's

```
1: x = f()  
2: if isinstance(x, int):  
3:     y = x  
4: else:  
5:     y = "y"  
6: #  
7: y
```

```
Use(f@line1) => Forward(...)  
Def(x@line1) => Expr(x, f())  
Use(x@line2) => Forward(Def(x@line1))  
Use(int@line2) => Forward(...)  
N(x@line2, line3) => N(Use(x@line2), IsInstance(int))  
N(x@line2, line4) => N(Use(x@line2), NotInstance(int))  
Use(x@line3) => Forward(N(x@line2, line3))  
Def(y@line3) => Expr(x)  
Def(y@line5) => Expr("y")  
Phi(y@line6) => Phi(Def(y@line3), Def(y@line5))  
Use(y@line7) => Forward(Phi(y@line6))
```



5. Answers (expressions)

- Solve the bindings that were created to produce types.
- Key \rightarrow Thunk<Type> mapping
 - Allows cycles
 - Use Var as a placeholder for unknown types, and unification.
- Lots and lots of code for each special type object in Python.



6. Interface (types)

- Types of exported symbols only.
- Subset of the Answers (less memory)



What are the types of x?

```
x = [1]
```

```
x = [1, "test"]
```

```
x: Literal[1] = [1]
```

```
x = []
```



The magic Var (unification, inference, loops)

```
x = list[?1]
add = Callable[[list[?2], ?2], None]
list[?2] = list[?1] ⇒ ?2 = ?1
?2 = Literal[1] ⇒ ?2 = int (generalise), ?1 = int
```

```
x = []
add(x, 1)
```

```
def add[T](a: list[T], b: T) -> None: ...
```



Eviction


- For most modules, we throw away AST, Bindings, Answers when done
- For files open in the IDE, we keep everything to support goto-def etc
- We always keep code around because the disk might change
- If you ask for the type of an export, we try the interface first, if not, the answers
 - Answers is required to deal with cycles
- Prioritise modules that are “nearly finished” to reduce memory

Eviction rules are simple because each module goes through sequential steps.



Incrementality (with cycles)

- A changes, what should we invalidate?
- Pyrefly says
 - Optimistically invalidate only A
 - Compute A using stale values of B/C
 - If the interface of A changes, invalidate B/C
 - Compute B/C using the last iteration of A
 - If A changes (since it changed before) invalidate the cycle
- Pros: Usually only one file invalidates
- Cons: Might compute a module more than once, and less parallelism





Performance

We want to check on every keystroke. We can check 1.85M lines/second.

Approximately Answers is 10x the cost of Bindings, which is 10x Exports.

We freely clone Type all the time. Should really have a heap...

Lots of profiling. Super easy to go quadratic.

With lots of threads, and Rust, the expensive things are:

- Thread communication
- Locks

Both are in the build system, which has been optimised a fair amount.



Transactions

An IDE does lots of things at once:

- Indexing (for find references)
- Checking a file that changed
- Answering queries

We have a transactional build system, with explicit commit. Never have to wait.

Can only have at most one mutable transaction at a time (so commit always succeeds)



Extensibility

- Model of features built in to the core
- Build systems like Buck/Bazel/Dune?
 - Buck integrated into the core, using queries and Bxl
 - Very open to further integrations
- Mypy plugins?
 - Supports Pydantic rules natively
 - Aiming to add some more special cases
 - Shape types one day?



Recursion

```
struct Thunk<T>;
```

- If this thread is already calculating this Thunk, create a Var.
- If not, solve the binding.
- When the calculation completes, bind the Var.



Open source

- We have gained much from open source!
 - Python itself
 - Python typing specification, plus existing checkers (Pyright, Mypy etc)
 - Ruff parser
 - Open source Python projects, e.g. PyTorch (now has Pyrefly in shadow)
- MIT license, <https://github.com/facebook/pyrefly>
- Delighted to accept pull requests, all issues are on issue tracker





The journey of autocomplete

```
display(3.142).fraction
```

```
from typing import *
from numbers import *
```

- Find the type of `display(3.142).fraction`
- First, find `display`
 - Might come from `typing` or `numbers`
 - Figure out the export table from each
 - Which might require a fixed-point of recursive `*` imports...



The journey of autocomplete (2)

```
@dataclass
class Number[T]:
    whole: T
    fraction: Final[T]

def display(x: float) -> Number[float]:
    whole = float(math.floor(x))
    fractional = x - whole
    return Number(whole, fractional)
```

- Interpret `@dataclass`
- Infer types for each variable
- Infer the return type
- Instantiate some generics
- Understand `Final`



The journey of autocomplete (3)

```
display(3.142).fraction.
```

```
as_integer_ratio
```

```
conjugate
```

```
fromhex
```

```
hex
```

```
imag
```

```
is_integer
```

```
real
```

```
__abs__
```

```
__add__
```


```
class float:  
    def __new__(cls, x = ...) -> Self  
    @classmethod  
    def fromhex(cls, x: str) -> Self  
    @property  
    def real(self) -> float  
    def conjugate(self) -> float  
    def __add__(self, x: float) -> float
```

- Now we know we have `float`
- Figure out what methods it has



Why not Pyrefly?

- It is an alpha - lots of known bugs
- You will probably find bugs, most of which we will fix
- But you will get a sticker (if you are here)





The team (+ over 100 contributors)





Questions?

pyrefly.org